

## Chapter 2.2.3 GKE networking

Pod Networking | Qwiklabs 3:52 PM 1/2/2021 InPrivate

### Pod Networking

Mark as Completed

A Pod is a group of containers with shared storage and networking

Pod

eth0  
10.4.0.2

Shared networking namespace (localhost 127.0.0.1)

nginx

legacy-app

tcp:80

tcp:8000

because both containers share the same networking name

Chat

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Increase Resources

Pod Networking | Qwiklabs 3:53 PM 1/2/2021 InPrivate

### Pod Networking

Mark as Completed

Your workload doesn't run in a single Pod

Google Cloud Platform

Production Namespace

Frontend Deployment

Backend Deployment

Staging Namespace

Frontend Deployment

Backend Deployment

Cloud Load Balancing

End Users

Ops & Quality

your workload is composed of many different applications

Chat

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Increase Resources

Pod Networking | Qwiklabs 3:53 PM 1/2/2021 InPrivate

## Pod Networking

**Pod-to-Pod communication on the same node**

using the node's vm nic the root network namespace

Chat

Pod Networking | Qwiklabs 3:54 PM 1/2/2021 InPrivate

## Pod Networking

**Nodes get Pod IP addresses from address ranges assigned to your Virtual Private Cloud**

Name	Zone	Internal IP	External IP
gke-standard-cluster-1-default-pool-echdfe5e-g3ar	us-central1-a	10.128.0.4 (node)	35.238.171.108
gke-standard-cluster-1-default-pool-echdfe5e-g3bm	us-central1-a	10.128.0.3 (node)	35.184.108.27
gke-standard-cluster-1-default-pool-echdfe5e-g3cg	us-central1-a	10.128.0.2 (node)	35.239.9.246

as kubernetes clusters compute engine instances

Chat

Pod Networking | Qwiklabs 3:55 PM 1/2/2021 InPrivate

## Pod Networking

Addressing the Pods

on gcp alias ips allow you to configure additional

Mark as Completed Chat

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Increase Resources

Pod Networking | Qwiklabs 3:55 PM 1/2/2021 InPrivate

## Pod Networking

Addressing the Pods

address so this address

Mark as Completed Chat

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Increase Resources

Pod Networking | Qwiklabs 3:56 PM 1/2/2021 InPrivate

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49563

Pod Networking

Addressing the Pods

Cluster-wide Service IP range (4,000)

Pod IP range (for the entire cluster /14)

Pod IPs /24

Pod IPs /24

Pod IPs /24

Node 1 NIC

Node 2 NIC

Node 3 NIC

VPC

google doesn't expect you to run 250 000 pods in a single

Mark as Completed

Chat

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Increase Resources

Pod Networking | Qwiklabs 3:56 PM 1/2/2021 InPrivate

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49563

Pod Networking

Pod-to-Pod communication

Pod 1 IP

Pod 6 IP

Pod IPs /24

Pod IPs /24

Pod IPs /24

Node 1 NIC

Node 2 NIC

Node 3 NIC

VPC

the nodes allocate a unique ip address from their assigned

Subnet details

Subnet range name	Secondary IP range
gke-standard-cluster-1-pods-0534accd	10.8.0.0/14
gke-standard-cluster-1-services-0534accd	10.12.0.0/20

Mark as Completed

Chat

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Increase Resources

Pod Networking | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49563

Pod Networking

Mark as Completed

Video Pod Networking

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Ingress Resource

## Communicating outside Google Cloud

The diagram illustrates a Kubernetes cluster consisting of three nodes. Each node contains multiple pods. Node 1 has Pod 1 (IP 10.8.1.2) and Pod 4 (IP 10.8.1.12). Node 2 has Pod 5 (IP 10.8.2.1). Node 3 has Pod 6 (IP 10.8.3.2). Each node is connected to a NIC (Network Interface Card). The nodes are interconnected by a VPC (Virtual Private Cloud). A curved arrow points from the cluster to a cloud icon labeled "Internet", indicating that traffic from the cluster is routed or peered inside the VPC to access the external Internet.

VPC

traffic from your cluster  
is routed or peered inside

Chat

Quiz: Pod networking | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49564

Quiz: Pod networking

Your score: 100% Passing score: 50%

Congratulations! You passed this assessment.

Retake

Quiz Quiz: Pod networking

Video Services

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Ingress Resource

Quiz Quiz: Ingress

✓ 1. Which statement is true about Kubernetes networking?

Each Pod in a node has a unique IP address, but IP addresses might be duplicated among nodes.

Each container in a cluster has a unique IP address.

Each Pod in a cluster has a unique IP address.

That is correct.

✓ 2. In GKE, what is the source of the IP addresses for Pods?

Loopback network addresses

Chat

Quiz: Pod networking | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49564

Quiz: Pod networking

That is correct.

That is correct.

2. In GKE, what is the source of the IP addresses for Pods?

- Loopback network addresses
- Arbitrary network addresses per cluster
- Address ranges assigned to your Virtual Private Cloud

Chat

Quiz: Pod networking

Video Services

Video Finding Services

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49565

Services

Mark as Completed

Pod Networking

Quiz: Pod networking

Video Services

Video Finding Services

Quiz: Services

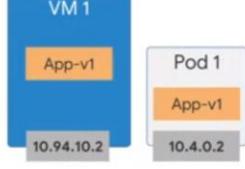
Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Pods versus VMs: Durability



different ways to find services in GKE  
virtual machines and pods

Chat

Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49565

Mark as Completed

Services

Pod Networking

- Quiz
- Quiz: Pod networking
- Video Services
- Video Finding Services
- Quiz Quiz: Services
- Video Service Types and Load Balancers
- Video How Load Balancers Work
- Quiz Quiz: Service types
- Video Ingress Resource

## Pods versus VMs: Durability

VM 1  
App-v2  
10.94.10.2

Pod 1  
App-v1  
10.4.0.2

Pod 2  
App-v2  
10.4.0.3

Terminate Deploy

through application updates and upgrades whereas pods are typically terminated

Chat

A woman in a blue Google t-shirt is speaking.

Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49565

Mark as Completed

Services

Pod Networking

- Quiz
- Quiz: Pod networking
- Video Services
- Video Finding Services
- Quiz Quiz: Services
- Video Service Types and Load Balancers
- Video How Load Balancers Work
- Quiz Quiz: Service types
- Video Ingress Resource

## Pods versus VMs: Durability

VM 1  
App-v2  
10.94.10.2

Pod 1  
App-v1  
10.4.0.2

Pod 2  
App-v2  
10.4.0.3

New Pod, New IP

gets a new IP address also if a part is rescheduled for

Chat

A woman in a blue Google t-shirt is speaking.

Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49565

Mark as Completed

Services

A Kubernetes Service creates endpoints

Frontend Pod → Service → Backend Pod

Endpoint

Frontend Pod → Service → Backend Pod

Frontend Pod → Service → Backend Pod

Backend Pod

Backend Pod

Backend Pod

Fortunately Kubernetes has an answer: services kubernetes service is an object that creates.

Chat

Pod Networking

Quiz

Quiz: Pod networking

Video Services

Video Finding Services

Quiz

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz

Quiz: Service types

Video Ingress Resource

Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49565

Mark as Completed

Services

A Service is issued a static Virtual IP address

Virtual IP → Frontend Pod → Service → Backend Pod

Frontend Pod

Service

Backend Pod

Backend Pod

Backend Pod

that the cluster reserves for services the virtual

Chat

Pod Networking

Quiz

Quiz: Pod networking

Video Services

Video Finding Services

Quiz

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz

Quiz: Service types

Video Ingress Resource

Finding Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49566

Finding Services

Mark as Completed

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

There are several ways to find a Service in GKE

Environment Variables

Kubernetes DNS

Istio

Chat

Finding Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49566

Finding Services

Mark as Completed

Video Finding Services

Quiz Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

Environment variables for Service discovery

Chat

Finding Services | Qwiklabs 4:16 PM 1/2/2021

[https://googlecourses.qwiklabs.com/course\\_sessions/96271/video/49566](https://googlecourses.qwiklabs.com/course_sessions/96271/video/49566)

Finding Services

Mark as Completed

Video Finding Services

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Example of environment variables for a Service

```
DEMO_SERVICE_HOST=10.70.0.11  
DEMO_SERVICE_PORT=6379  
DEMO_PORT=tcp://10.70.0.11:6379  
DEMO_PORT_6379_TCP=tcp://10.70.0.11:6379  
DEMO_PORT_6379_TCP_PROTO=tcp  
DEMO_PORT_6379_TCP_PORT=6379  
DEMO_PORT_6379_TCP_ADDR=10.70.0.11
```

demo.my-project

\_http.\_tcp.demo.my-project

Chat

Finding Services | Qwiklabs 4:16 PM 1/2/2021

[https://googlecourses.qwiklabs.com/course\\_sessions/96271/video/49566](https://googlecourses.qwiklabs.com/course_sessions/96271/video/49566)

Finding Services

Mark as Completed

Video Finding Services

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Service discovery through Kubernetes DNS

Namespaces: demo

Service "lab"

10.70.0.11

Pod

App

Cluster Master

api-server

kube-dns

lab ?

Chat

Finding Services

Mark as Completed

Services are assigned one A (Address) record and one SRV (Service) record

Kubernetes DNS

Service “lab”

Namespace: “demo”

record type: A (Address)  
hostname: lab  
namespace: demo  
FQDN: lab.demo.svc.cluster.local  
IP address: 10.70.0.10

Chat

Video Finding Services

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Finding Services

Mark as Completed

Services are assigned one A (Address) record and one SRV (Service) record

Kubernetes DNS

Service “lab”

Namespace: “demo”

record type: SRV (Service)  
hostname: lab  
namespace: demo  
FQDN: \_http.\_tcp.lab.demo.svc.cluster.local  
Value: 80

Chat

Video Finding Services

Quiz: Services

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Finding Services | Qwiklabs 4:19 PM 1/2/2021 InPrivate

[Video Finding Services](#)

[Quiz Quiz: Services](#)

[Video Service Types and Load Balancers](#)

[Video How Load Balancers Work](#)

[Quiz Quiz: Service types](#)

[Video Ingress Resource](#)

[Quiz Quiz: Ingress](#)

[Video Container-Native Load Balancing](#)

[Mark as Completed](#)

## Service discovery through Istio

Open, platform-independent service mesh

Adds visibility and control to a microservices architecture

Available as a plug-in in GKE<sup>beta</sup>



A woman in a blue Google t-shirt stands to the right of the slide.

Chat

Finding Services | Qwiklabs 4:20 PM 1/2/2021 InPrivate

[Video Finding Services](#)

[Quiz Quiz: Services](#)

[Video Service Types and Load Balancers](#)

[Video How Load Balancers Work](#)

[Quiz Quiz: Service types](#)

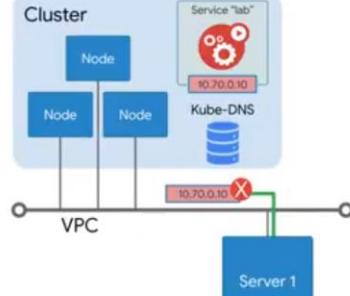
[Video Ingress Resource](#)

[Quiz Quiz: Ingress](#)

[Video Container-Native Load Balancing](#)

[Mark as Completed](#)

## Service discovery outside the cluster



A woman in a blue Google t-shirt stands to the right of the slide.

Chat

Quiz: Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49567

InPrivate

## Quiz: Services

Congratulations! You passed this assessment.

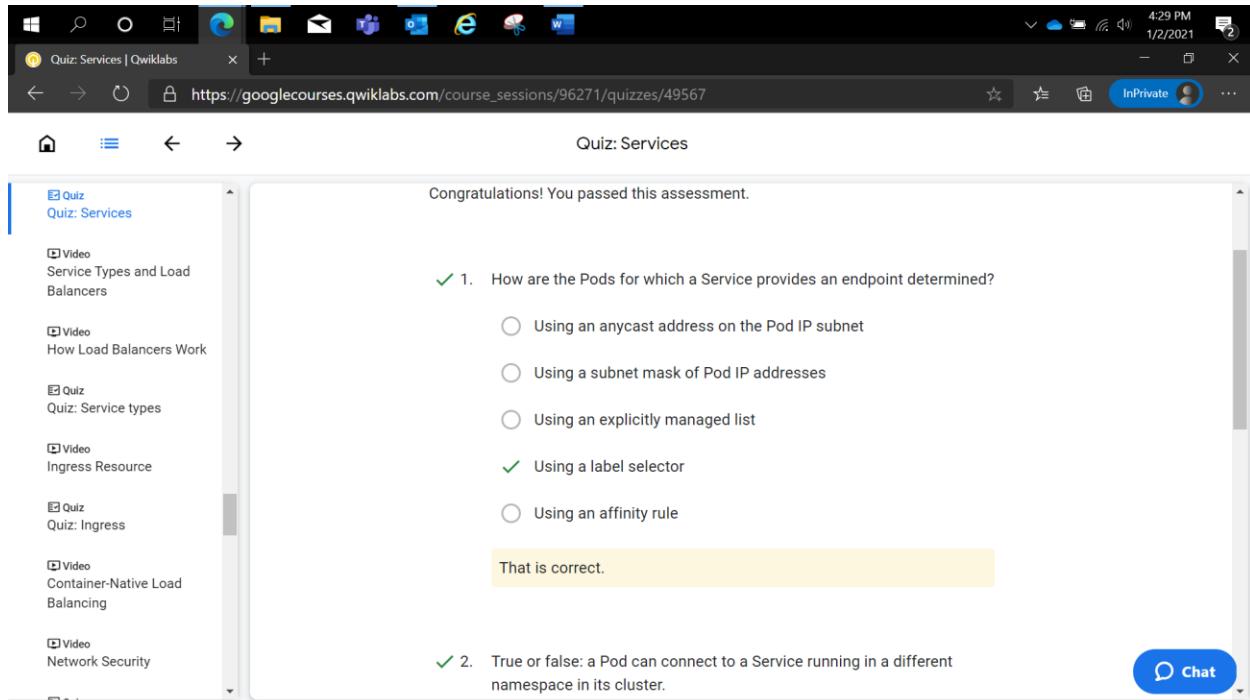
1. How are the Pods for which a Service provides an endpoint determined?

- Using an anycast address on the Pod IP subnet
- Using a subnet mask of Pod IP addresses
- Using an explicitly managed list
- Using a label selector
- Using an affinity rule

That is correct.

2. True or false: a Pod can connect to a Service running in a different namespace in its cluster.

Chat



Quiz: Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49567

InPrivate

## Quiz: Services

That is correct.

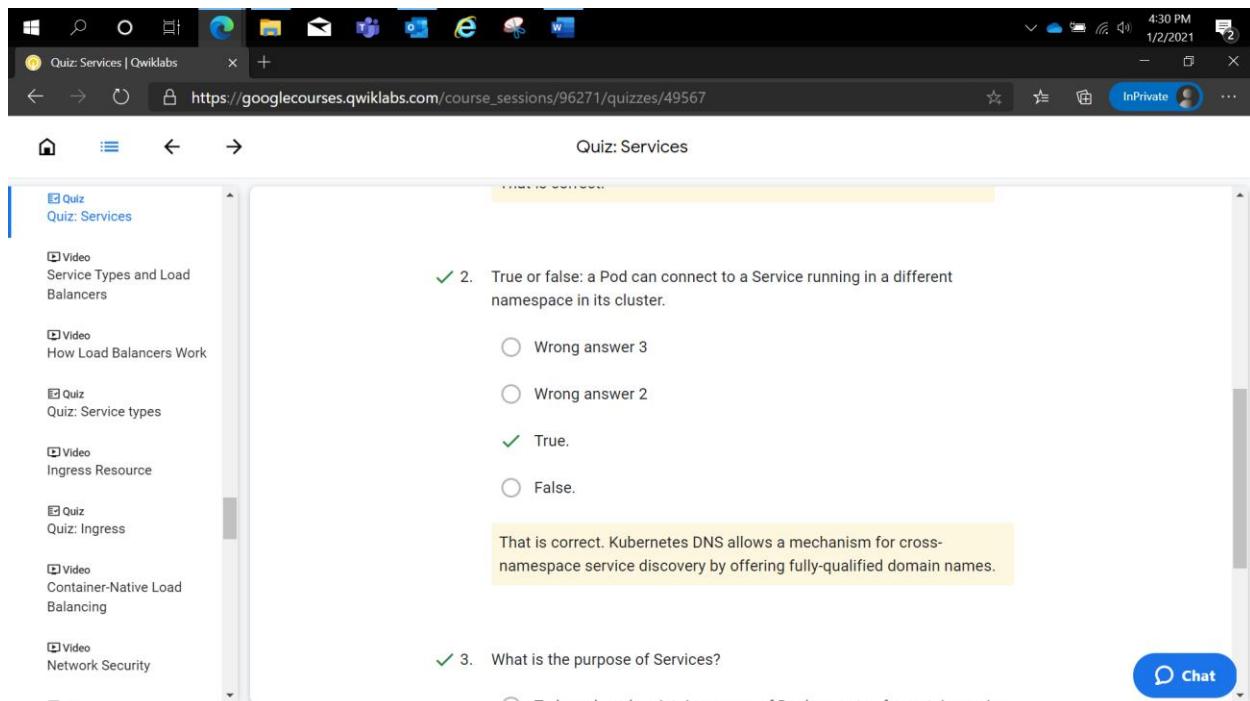
2. True or false: a Pod can connect to a Service running in a different namespace in its cluster.

- Wrong answer 3
- Wrong answer 2
- True.
- False.

That is correct. Kubernetes DNS allows a mechanism for cross-namespace service discovery by offering fully-qualified domain names.

3. What is the purpose of Services?

Chat



Quiz: Services | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49567

Quiz: Services

That is correct. Kubernetes DNS allows a mechanism for cross-namespace service discovery by offering fully-qualified domain names.

3. What is the purpose of Services?

- To launch and maintain a group of Deployments of a certain version.
- To launch and maintain a group of Pods of a certain version.
- To provide a static, load-balanced front end for transient Pods.
- To provide a static, load-balanced front end for kube-apiserver.

That is correct.

Chat

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Service Types and Load Balancers | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49568

Service Types and Load Balancers

Mark as Completed

A ClusterIP Service has a static IP address

the cluster ip service a kubernetes cluster ip service has a static ip

Chat

Video Service Types and Load Balancers

Video How Load Balancers Work

Quiz Quiz: Service types

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz Quiz: Network security

Video

Service Types and Load Balancers

Setting up a ClusterIP Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: Backend
  ports:
    - protocol: TCP
      port: 3306
      targetPort: 6000
```

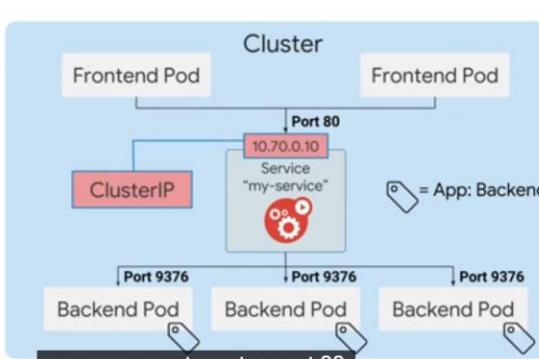
step-by-step process of setting up a cluster



Chat

Service Types and Load Balancers

How does the ClusterIP Service work?



Cluster

Frontend Pod → Port 80 → 10.70.0.10 → Service "my-service" → Port 9376 → Backend Pod

Frontend Pod → Port 80 → 10.70.0.10 → Service "my-service" → Port 9376 → Backend Pod

Frontend Pod → Port 80 → 10.70.0.10 → Service "my-service" → Port 9376 → Backend Pod

= App: Backend

answer requests on tcp port 80.  
kubernetes manages

Chat

Service Types and Load Balancers

Mark as Completed

NodePort Service

of the back end pods  
on port 9376

default NodePort Range (30000-32767)

Chat

Service Types and Load Balancers

Mark as Completed

Creating a NodePort Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: NodePort
  selector:
    app: Backend
  ports:
    - protocol: TCP
      nodePort: 30100
      port: 80
      targetPort: 9376
```

automatically allocated from the range  
30 000 to 32700

Chat

How Load Balancers Work | + https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49569

## How Load Balancers Work

Mark as Completed

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz: Network security

Video Lab Intro

Hands-On Lab <https://www.youtube.com/watch?v=l8Qwl2kDl34>

### LoadBalancer Service

A diagram titled "LoadBalancer Service" showing the architecture. A "Client" connects to a "LoadBalancer Service" (represented by a triangle icon). The "LoadBalancer Service" then connects to three "Node"s. Each "Node" connects to an "External Pod Cluster". The diagram shows three "External Pod Cluster"s, each with a port mapping from TCP: 80 on the "Node" to TCP: 9376 on the "External Pod".

Chat

Google

How Load Balancers Work | + https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49569

## How Load Balancers Work

Mark as Completed

Video How Load Balancers Work

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz: Network security

Video Lab Intro

Hands-On Lab [Configuring Google](#)

### Creating a LoadBalancer Service

```
[...]
spec:
  type: LoadBalancer
  selector:
    app: Backend
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

Chat

Google

How Load Balancers Work | 4:46 PM 1/2/2021

## Service types summary

The diagram illustrates three service types:

- ClusterIP:** A ClusterIP service is represented by a red circle with a port number (TCP:30100). It receives traffic on Port 80 from a Client and routes it to two Frontend Pods (TCP:30100) and three Backend Pods (TCP:9376). The Backend Pods are connected to a ClusterIP (TCP:30100).
- NodePort:** A NodePort service is represented by a blue rectangle labeled "NodePort Service". It receives traffic on Port 80 from a Client and routes it to three Nodes (TCP:30100). Each Node is connected to an External Pod Cluster (TCP:9376).
- LoadBalancer:** A LoadBalancer service is represented by a green rectangle labeled "LoadBalancer Service". It receives traffic on Port 80 from a Client and routes it to three Nodes (TCP:30100). Each Node is connected to an External Pod Cluster (TCP:9376).

Quiz: Service types

Video Ingress Resource

Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Mark as Completed

Chat

Quiz: Service types | Qwiklabs | 4:47 PM 1/2/2021

## Congratulations! You passed this assessment.

✓ 1. What is the difference between a Service of type ClusterIP and one of type NodePort?

In addition to exposing the service at a particular cluster IP address, a NodePort Service also exposes the service on every node at a particular port number.

In addition to exposing the service at a particular cluster IP address, a NodePort Service also exposes the service on a Google Cloud Network Load Balancer.

That is correct.

✓ 2. What is the main use of a Service of type LoadBalancer?

Exposing services outside a cluster

Internal communication within a cluster

Chat

Quiz: Service types | Qwiklabs 4:47 PM 1/2/2021

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49570 InPrivate

### Quiz: Service types

Quiz: Service types

2. What is the main use of a Service of type LoadBalancer?

Exposing services outside a cluster

Internal communication within a cluster

That is correct.

3. What is the main use of a Service of type ClusterIP?

Internal communication within a cluster

Exposing services outside a cluster

That is correct.

Chat

Ingress Resource | Qwiklabs 4:48 PM 1/2/2021

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571 InPrivate

### Ingress Resource

Mark as Completed

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

The Ingress resource, a service for Services

Traffic

The diagram illustrates the role of an Ingress resource in a Kubernetes cluster. External traffic enters through an Ingress object, which then routes requests to different services within the cluster. One service handles traffic for 'sales.mydomain.com', another for 'mydomain.com/video', and a third for 'Other' traffic. Each service is connected to multiple pods, representing the actual application instances.

Kubernetes cluster

Service for Services Ingress is not a service or

Chat

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

Ingress Resource

Mark as Completed

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Google Kubernetes Engine specifics

GKE deploys a Google Cloud HTTP(S) Load Balancer

Service considerations

Ingress can direct traffic to:

- NodePort Services
- LoadBalancer Services

LoadBalancer Services still have the double-hop problem

constructs and can deliver traffic to either node port services

Settings Chat

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

Ingress Resource

Mark as Completed

Video Ingress Resource

Quiz Quiz: Ingress

Video Container-Native Load Balancing

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Creating an Ingress

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: test-ingress
spec:
  backend:
    serviceName: demo
    servicePort: 80
```

Traffic

Ingress: my-Ingress

Service demo:80

Kubernetes cluster

here in the object

Chat

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

## Ingress manifest example 2

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-Ingress
spec:
  rules:
    - host: demo1.example.com
      http:
        paths:
          - path: /demolexamplepath
            backend:
              serviceName: demo
              servicePort: 80
```

Traffic

Ingress: my-Ingress  
demo1.example.com/demo1examplepath

Service demo:80

Kubernetes cluster

the specifications  
there are

Chat

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

## Ingress manifest example 3

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-Ingress
spec:
  rules:
    - host: demo.example.com
      http:
        paths:
          - path: /demopath
            backend:
              serviceName: demo1
              servicePort: 80
```

```
- host: lab.user.com
  http:
    paths:
      - path: /labpath
        backend:
          serviceName: lab1
          servicePort: 80
```

supports  
multiple host names for the same

Chat

Ingress Resource

### Ingress manifest example 3

The diagram illustrates an Ingress resource named "my-Ingress" with two rules. The first rule maps the host "demo.example.com/demopath" to the service "demo1:80", which is backed by three pods. The second rule maps the host "lab.user.com/labpath" to the service "lab1:80", which is also backed by three pods. Arrows indicate the flow of traffic from the external world through the Ingress to the respective services and then to the pods.

**Traffic**

**Ingress: my-Ingress**

demo.example.com/demopath      lab.user.com/labpath

**Kubernetes cluster**

the traffic will be redirected from the http or https

Ingress Resource

### Ingress manifest example 4

```

apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-Ingress
spec:
  rules:
    - host: demo.example.com
      http:
        paths:
          - path: /demopath
            backend:
              serviceName: demo1
              servicePort: 80

```

```

          -path: /demo2path
          backend:
            serviceName: demo2
            servicePort: 80

```

to the service named demo1 on port 80.  
this example considers rules

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

## Ingress manifest example 4

The diagram illustrates an Ingress configuration named "my-Ingress". It receives "Traffic" from the outside world. Two URL paths are defined: "demo.example.com/demo1path" and "demo.example.com/demo2path". These paths are mapped to two different services: "demo1:80" and "demo2:80". Each service then routes traffic to a group of three "Pod"s within a "Kubernetes cluster".

url path here the traffic from demo.example.com

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

## Updating an Ingress

\$ kubectl edit ingress [NAME]

\$ kubectl replace -f [FILE]

replace command which replaces the ingress object

Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

Ingress Resource

Mark as Completed

Quiz: Service types

Video: Ingress Resource

Quiz: Ingress

Video: Container-Native Load Balancing

Video: Network Security

Quiz: Network security

Video: Lab Intro

Hands-On Lab: Configuring Google Kubernetes Engine (GKE) Networking

Ingress natively supports many Google Cloud services

IAP

Cloud Armor

Cloud CDN

cloud content delivery network allows you to bring

Chat



Ingress Resource | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49571

Ingress Resource

Mark as Completed

Quiz: Service types

Video: Ingress Resource

Quiz: Ingress

Video: Container-Native Load Balancing

Video: Network Security

Quiz: Network security

Video: Lab Intro

Hands-On Lab: Configuring Google Kubernetes Engine (GKE) Networking

Additional Ingress features

- TLS termination
- Multiple SSL certificates
- HTTP/2 and gRPC
- Multi-cluster and multi-region support

within your cluster lastly with multi-cluster

Chat

Quiz: Ingress | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49572

Quiz: Ingress

Your score: 100% Passing score: 50%

Congratulations! You passed this assessment.

**1.** How does an Ingress resource decide how to route incoming requests?

- By the host name and resource path requested
- By the IP address requested
- By the host name requested
- By the port number requested

That is correct.

**2.** What is the main purpose of an Ingress resource?

- By the port number requested

That is correct.

Retake

Chat

Quiz: Ingress

- Quiz Quiz: Ingress
- Video Container-Native Load Balancing
- Video Network Security
- Quiz Quiz: Network security
- Video Lab Intro
- Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking
- Video Lab Intro
- Hands-On Lab Creating Services and Ingress Resources

Quiz: Ingress | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49572

Quiz: Ingress

By the port number requested

That is correct.

**2.** What is the main purpose of an Ingress resource?

- To provide external access to one or more Services
- To provide internal access to one or more Services
- To provide external access to exactly one Service

That is correct.

Chat

Quiz: Ingress

- Quiz Quiz: Ingress
- Video Container-Native Load Balancing
- Video Network Security
- Quiz Quiz: Network security
- Video Lab Intro
- Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking
- Video Lab Intro
- Hands-On Lab Creating Services and Ingress Resources

Container-Native Load Balancing

The double-hop dilemma

Cluster

Pod 1 Pod 2

Pod 3 Pod 4

Pod 5 Pod 6

iptables

Node 1 Node 2 Node 3

NIC

VPC

Client

Container-Native Load Balancing

Add the Local value for the external-traffic policy

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: LoadBalancer
  externalTrafficPolicy: Local
  selector:
    app: external
  ports:
    - protocol: TCP
      port: 80
      targetPort: 9376
```

Chat

Container-Native Load Balancing

Container-native load balancer: The details

- Eliminates the “Local” external-traffic policy workaround
- Leverages a Google Cloud HTTP(S) load balancer
- Traffic is directed to the Pods directly instead of to the nodes
- Uses a data model called Network Endpoint Groups



Chat

Video Container-Native Load Balancing

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

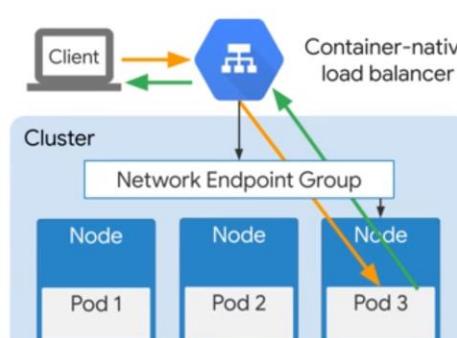
Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Networking

Container-Native Load Balancing

Container-native load balancer in action



```
graph LR; Client --> LN[Container-native load balancer]; LN --> NEG[Network Endpoint Group]; NEG --> Node1[Node Pod 1]; NEG --> Node2[Node Pod 2]; NEG --> Node3[Node Pod 3]
```

Client

Container-native load balancer

Cluster

Network Endpoint Group

Node

Pod 1

Pod 2

Pod 3

Chat

Video Container-Native Load Balancing

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Networking

Container-Native Load Balancing

Container-Native Load Balancing

Benefits

- 1 Traffic is appropriately directed
- 2 Support for load balancer features
- 3 Increased visibility
- 4 Improved network performance
- 5 Support for other Google Cloud networking services



Chat

Network Security | Qwiklabs

Network Security

Network policy

A Pod-level firewall restricting access to other Pods and Services

Must be enabled:

- Requires at least 2 nodes of n1-standard-1 or higher
- Requires nodes to be recreated



Chat

Network Security | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49574

Network Security

Mark as Completed

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Engine Networking

Video Summary

## Enabling a network policy

Enable a network policy for a new cluster

```
gCloud container clusters create [NAME] \
--enable-network-policy
```

Enable a network policy for an existing cluster

```
gCloud container clusters update [NAME] \
--update-addons-NetworkPolicy=ENABLED
gCloud container clusters update [NAME] \
--enable-network-policy
```



Chat

Network Security | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49574

Network Security

Mark as Completed

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Engine Networking

Video Summary

## Writing a network policy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: demo-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: demo-app
  policyTypes:
    - Ingress
    - Egress
```

```
ingress:
  - from:
    - ipBlock:
        cidr: 172.17.0.0/16
      except:
        - 172.17.1.0/24
    - namespaceSelector:
        matchLabels:
          project: myproject
  - podSelector:
      matchLabels:
        role: frontend
  ports:
    - protocol: TCP
      port: 6379
```

Chat

Network Security | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49574

Network Security

Mark as Completed

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Engine Networking

Video Summary

## Writing a network policy

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: demo-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: demo-app
  policyTypes:
    - Ingress
    - Egress
```

```
egress:
  - to:
    - ipBlock:
      cidr: 10.0.0.0/24
  ports:
    - protocol: TCP
      port: 5978
```

Chat

Network Security | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49574

Network Security

Mark as Completed

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Engine Networking

Video Summary

## Disabling a network policy

```
Disable a network policy for a cluster
gcloud container clusters create [NAME] \
  --no-enable-network-policy
```



Chat

Network Security | Qwiklabs

https://googlecourses.qwiklabs.com/course\_sessions/96271/video/49574

Network Security

Mark as Completed

Video Network Security

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Engine Networking

Video summary

## Network policy defaults

```
metadata:
  name: default-deny
spec:
  podSelector: {}
  policyTypes:
    - Ingress
```

```
metadata:
  name: allow-all
spec:
  podSelector: {}
  policyTypes:
    - Ingress
    - Egress
```

```
metadata:
  name: default-deny
spec:
  podSelector: {}
  policyTypes:
    - Egress
```

```
metadata:
  name: allow-all
spec:
  podSelector: {}
  policyTypes:
    - Egress
  egress:
    - {}
```

Chat

Quiz: Network security | Qwiklab

https://googlecourses.qwiklabs.com/course\_sessions/96271/quizzes/49575

Quiz: Network security

Your score: 100% Passing score: 50%

Congratulations! You passed this assessment.

Retake

Quiz Quiz: Network security

Video Lab Intro

Hands-On Lab Configuring Google Kubernetes Engine (GKE) Networking

Video Lab Intro

Hands-On Lab Creating Services and Ingress Resources

Quiz Quiz: Google Kubernetes Engine Networking

Video summary

Persistent Data and

✓ 1. What is the purpose of enabling network policies in a Kubernetes cluster?

To restrict network access from outside a cluster to Pods inside the cluster

To restrict network access from a Pod to other Pods and Services inside the cluster

To restrict network access from outside a cluster to Services inside the cluster

That is correct.

✓ 2. What are the requirements for enabling network security policies? Choose all answers that are correct (3 correct answers).

Chat

Quiz: Network security

✓ 2. What are the requirements for enabling network security policies? Choose all answers that are correct (3 correct answers).

✓ All nodes must be of machine type n1-standard-1 or higher.  
That is correct.

You must install a software firewall module on each node.

✓ You must have at least 2 nodes.  
That is correct.

✓ You must restart your nodes.  
That is correct.

You must have enabled container-native load balancing.

Kubernetes Engine

Kubernetes clusters

Name	Location	Cluster size	Total cores	Total memory	Notifications	Labels
private-cluster	us-central1-a	2	4 vCPUs	8.00 GB		

CLOUD SHELL

```
version: 1.16.15-gke.4901
privateClusterConfig:
enablePrivateNodes: true
masterIpv4CidrBlock: 172.16.0.0/28
peeringName: gke-n338e33900bffc258d4d-103e-b1b0-peer
privateEndpoint: 172.16.0.2
publicEndpoint: 34.72.209.160
releaseChannel: {}
selfLink: https://container.googleapis.com/v1/projects/qwiklabs-gcp-03-2ac151202de9/zones/us-central1-a/clusters/private-cluster
servicesIpv4Cidr: 10.36.0.0/20
shieldedNodes: {}
status: RUNNING
subnetwork: default
zone: us-central1-a
student_03_c10d2d4f55e9@cloudshell:~ (qwiklabs-gcp-03-2ac151202de9)$
```

# Configuring Google Kubernetes Engine (GKE) Networking

1 hourFree

Rate Lab

## Overview

In this lab, you will create a private cluster, add an authorized network for API access to it, and then configure a network policy for Pod security.

## Objectives

In this lab, you learn how to perform the following tasks:

- Create and test a private cluster
- Configure a cluster for authorized network master access
- Configure a Cluster network policy

## Access Qwiklabs

For each lab, you get a new GCP project and set of resources for a fixed time at no cost.

1. Make sure you signed into Qwiklabs using an **incognito window**.
2. Note the lab's access time (for example, **02:00:00**) and make sure you can finish in that time block.

There is no pause feature. You can restart if needed, but you have to start at the beginning.

**START LAB**

3. When ready, click .

4. Note your lab credentials. You will use them to sign in to Cloud Platform Console.

[Open Google Console](#)

**Caution:** When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked. [Learn more.](#)

Username  
 [!\[\]\(4118c652257c5ffeac89ab3cf4f37e96\_img.jpg\)](#)

Password  
 [!\[\]\(e46628e188a6e3e8ee8033782765120f\_img.jpg\)](#)

GCP Project ID  
 [!\[\]\(1c3366fded475ba3fa51be3722595f6f\_img.jpg\)](#)

New to labs? [View our introductory video!](#)

5. Click **Open Google Console**.
6. Click **Use another account** and copy/paste credentials for **this** lab into the prompts.

If you use other credentials, you'll get errors or **incur charges**.

7. Accept the terms and skip the recovery resource page.

Do not click **End Lab** unless you are finished with the lab or want to restart it.

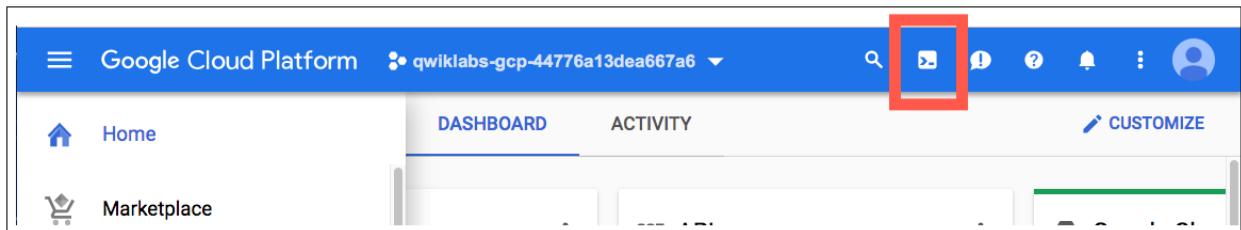
This clears your work and removes the project.

After you complete the initial sign-in steps, the project dashboard appears.

Activate Google Cloud Shell

Google Cloud Shell is a virtual machine that is loaded with development tools. It offers a persistent 5GB home directory and runs on the Google Cloud. Google Cloud Shell provides command-line access to your GCP resources.

1. In GCP console, on the top right toolbar, click the Open Cloud Shell button.



2. Click **Continue**.

A screenshot of a "Cloud Shell" setup page. At the top left is a "Cloud Shell" icon. Below it is a descriptive text block: "Google Cloud Shell provides you with command-line access to your cloud resources directly from your browser. You can easily manage your projects and resources without having to install the Google Cloud SDK or other tools on your system. [Learn more.](#)". At the bottom is a large blue "Continue" button.

It takes a few moments to provision and connect to the environment. When you are connected, you are already authenticated, and the project is set to your *PROJECT\_ID*. For example:

A screenshot of a Cloud Shell terminal window. The title bar says "...abs-gcp-44776a13dea667a6". The terminal output shows: "Welcome to Cloud Shell! Type "help" to get started.", "Your Cloud Platform project in this session is set to **qwiklabs-gcp-44776a13dea667a6**.", "Use "gcloud config set project [PROJECT\_ID]" to change to a different project.", and "google1623327\_student@cloudshell:~ (qwiklabs-gcp-44776a13dea667a6) \$". A large red arrow points from the text above to the "qwiklabs-gcp-44776a13dea667a6" part of the terminal output.

**gcloud** is the command-line tool for Google Cloud Platform. It comes pre-installed on Cloud Shell and supports tab-completion.

You can list the active account name with this command:

```
gcloud auth list  
content_copy
```

Output:

```
Credentialed accounts:  
- <myaccount>@<mydomain>.com (active) content_copy
```

Example output:

```
Credentialed accounts:  
- google1623327_student@qwiklabs.netcontent_copy
```

You can list the project ID with this command:

```
gcloud config list project  
content_copy
```

Output:

```
[core]  
project = <project_ID>content_copy
```

Example output:

```
[core]  
project = qwiklabs-gcp-44776a13dea667a6content_copy
```

Full documentation of **gcloud** is available on [Google Cloud gcloud Overview](#).

## Task 1. Create a private cluster

In this task, you create a private cluster, consider the options for how private to make it, and then compare your private cluster to your original cluster.

In a private cluster, the nodes have internal RFC 1918 IP addresses only, which ensures that their workloads are isolated from the public Internet. The nodes in a non-private cluster have external IP addresses, potentially allowing traffic to and from the internet.

## Set up a private cluster

1. On the **Navigation menu** (≡), click **Kubernetes Engine > Clusters**.
2. Click **Create cluster**.
3. Name the cluster `private-cluster`.
4. Select `us-central1-a` as the zone.
5. Click on **default-pool** under **NODE POOLS** section and then enter `2` in **Number of nodes** section.
6. Click on **Networking** section, select **Enable VPC-native traffic routing (uses alias IP)**.
7. In the **Networking** section, select **Private cluster** and select **Access master using its external IP address**.
8. For **Master IP Range**, enter `172.16.0.0/28`.

**Note:** This step assumes that you don't have this range assigned to a subnet. Check VPC Networks in the Google Cloud Console, and select a different range if necessary. Behind the scenes, Google Cloud uses VPC peering to connect the VPC of the cluster master with your default VPC network.

9. Deselect **Enable master authorized networks**.

This setting allows you to define the range of addresses that can access the cluster master externally. When this checkbox is not selected, you can access `kubectl` only from within the Google Cloud network. In this lab, you will only access `kubectl` through the Google Cloud network but you will modify this setting later.

10. Click **Create**.

**Note:** You need to wait a few minutes for the cluster deployment to complete.

## Inspect your cluster

1. In the Cloud Shell, enter the following command to review the details of your new cluster:

```
gcloud container clusters describe private-cluster --region us-central1-a  
content_copy
```

The following values appear only under the private cluster:

- `privateEndpoint`, an internal IP address. Nodes use this internal IP address to communicate with the cluster master.

- `publicEndpoint`, an external IP address. External services and administrators can use the external IP address to communicate with the cluster master.  
You have several options to lock down your cluster to varying degrees:

- The whole cluster can have external access.
- The whole cluster can be private.
- The nodes can be private while the cluster master is public, and you can limit which external networks are authorized to access the cluster master.  
Without public IP addresses, code running on the nodes can't access the public internet unless you configure a NAT gateway such as Cloud NAT.

You might use private clusters to provide services such as internal APIs that are meant only to be accessed by resources inside your network. For example, the resources might be private tools that only your company uses. Or they might be backend services accessed by your frontend services, and perhaps only those frontend services are accessed directly by external customers or users. In such cases, private clusters are a good way to reduce the surface area of attack for your application.

Click *Check my progress* to verify the objective.

Create a private cluster

Check my progress

## Task 2. Add an authorized network for cluster master access

After cluster creation, you might want to issue commands to your cluster from outside Google Cloud. For example, you might decide that only your corporate network should issue commands to your cluster master. Unfortunately, you didn't specify the authorized network on cluster creation.

In this task, you add an authorized network for cluster master access.

In this task, you make the Kubernetes master API accessible to a specific range of network addresses. In a real-world use of GKE, this connection would be used by IT staff and automated processes, not end-users.

1. In the Google Cloud Console **Navigation menu** (≡), click **Kubernetes Engine > Clusters**.
2. Click **private-cluster** to open the Clusters details page.
3. Click on **Edit** button.
4. In the **Master authorized networks** drop-down list, select **Enabled**.
5. Click **Add authorized network**.
6. For **Name**, type the name for the network, use `Corporate`.
7. For **Network**, type a CIDR range that you want to grant whitelisted access to your cluster master. As an example, you can use `192.168.1.0/24`.
8. Click **Done**.

Multiple networks can be added here if necessary, but no more than 50 CIDR ranges.

Outside this lab environment, a practical example might be to whitelist only the public, outside address of your corporate firewall. For example, if your corporate firewall's IP address were `8.8.8.14`, you could whitelist access to `8.8.8.14/32`.

9. Click **Save** at the bottom of the menu.

Click *Check my progress* to verify the objective.

Add an authorized network for cluster master access

Check my progress

## Task 3. Create a cluster network policy

In this task, you create a cluster network policy to restrict communication between the Pods. A zero-trust zone is important to prevent lateral attacks within the cluster when an intruder compromises one of the Pods.

# Create a GKE cluster

1. In Cloud Shell, type the following command to set the environment variable for the zone and cluster name.

```
export my_zone=us-central1-a  
export my_cluster=standard-cluster-1  
content_copy
```

2. Configure kubectl tab completion in Cloud Shell.

```
source <(kubectl completion bash)  
content_copy
```

3. In Cloud Shell, type the following command to create a Kubernetes cluster. Note that this command adds the additional flag `--enable-network-policy` to the parameters you have used in previous labs. This flag allows this cluster to use cluster network policies.

```
gcloud container clusters create $my_cluster --num-nodes 3 --enable-ip-alias  
--zone $my_zone --enable-network-policy  
content_copy
```

4. In Cloud Shell, configure access to your cluster for the kubectl command-line tool, using the following command:

```
gcloud container clusters get-credentials $my_cluster --zone $my_zone  
content_copy
```

5. Run a simple web server application with the label `app=hello`, and expose the web application internally in the cluster.

```
kubectl run hello-web --labels app=hello \  
--image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose  
content_copy
```

6. In Cloud Shell enter the following command to clone the repository to the lab Cloud Shell.

```
git clone https://github.com/GoogleCloudPlatform/training-data-analyst  
content_copy
```

7. Create a soft link as a shortcut to the working directory.

```
ln -s ~/training-data-analyst/courses/ak8s/v1.1 ~/ak8s  
content_copy
```

8. Change to the directory that contains the sample files for this lab.

```
cd ~/ak8s/GKE_Networks/  
content_copy
```

## Restrict incoming traffic to Pods

A sample NetworkPolicy manifest file called `hello-allow-from-foo.yaml` has been provided for you. This manifest file defines an ingress policy that allows access to Pods labeled `app: hello` from Pods labeled `app: foo`.

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: hello-allow-from-foo
spec:
  policyTypes:
    - Ingress
  podSelector:
    matchLabels:
      app: hello
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: foo
content_copy
```

1. Create an ingress policy.

```
kubectl apply -f hello-allow-from-foo.yaml
content_copy
```

2. Verify that the policy was created.

```
kubectl get networkpolicy
```

```
content_copy
```

### Output (Do not copy)

NAME	POD-SELECTOR	AGE
hello-allow-from-foo	app=hello	7s

```
content_copy
```

## Validate the ingress policy

1. Run a temporary Pod called `test-1` with the label `app=foo` and get a shell in the Pod.

```
kubectl run test-1 --labels app=foo --image=alpine --restart=Never --rm --
stdin --tty
content_copy
```

**Note:** The kubectl switches used here in conjunction with the run command are important to note.

--stdin ( alternatively -i ) creates an interactive session attached to STDIN on the container.

--tty ( alternatively -t ) allocates a TTY for each container in the pod.

--rm instructs Kubernetes to treat this as a temporary Pod that will be removed as soon as it completes its startup task. As this is an interactive session it will be removed as soon as the user exits the session.

--label ( alternatively -l ) adds a set of labels to the pod.

--restart defines the restart policy for the Pod

2. Make a request to the hello-web:8080 endpoint to verify that the incoming traffic is allowed.

```
wget -qO- --timeout=2 http://hello-web:8080  
content_copy
```

#### Output (Do not copy)

```
If you don't see a command prompt, try pressing enter.  
/ # wget -qO- --timeout=2 http://hello-web:8080  
Hello, world!  
Version: 1.0.0  
Hostname: hello-web-8b44b849-k96lh  
/ #  
content_copy
```

3. Type **exit** and press **ENTER** to leave the shell.

4. Now you will run a different Pod using the same Pod name but using a label, `app=other`, that does not match the `podSelector` in the active network policy. This Pod should not have the ability to access the hello-web application.

```
kubectl run test-1 --labels app=other --image=alpine --restart=Never --rm --  
stdin --tty  
content_copy
```

5. Make a request to the hello-web:8080 endpoint to verify that the incoming traffic is not allowed.

```
wget -qO- --timeout=2 http://hello-web:8080  
content_copy
```

The request times out.

## Output (Do not copy)

```
If you don't see a command prompt, try pressing enter.  
/ # wget -qO- --timeout=2 http://hello-web:8080  
wget: download timed out  
/ #  
content_copy
```

6. Type **exit** and press ENTER to leave the shell.

## Restrict outgoing traffic from the Pods

You can restrict outgoing (egress) traffic as you do incoming traffic. However, in order to query internal hostnames (such as hello-web) or external hostnames (such as www.example.com), you must allow DNS resolution in your egress network policies. DNS traffic occurs on port 53, using TCP and UDP protocols.

The NetworkPolicy manifest file `foo-allow-to-hello.yaml` file that has been provided for you defines a policy that permits Pods with the label `app: foo` to communicate with Pods labeled `app: hello` on any port number, and allows the Pods labeled `app: foo` to communicate to any computer on UDP port 53, which is used for DNS resolution. Without the DNS port open, you will not be able to resolve the hostnames.

```
kind: NetworkPolicy  
apiVersion: networking.k8s.io/v1  
metadata:  
  name: foo-allow-to-hello  
spec:  
  policyTypes:  
    - Egress  
  podSelector:  
    matchLabels:  
      app: foo  
  egress:  
    - to:  
      - podSelector:  
          matchLabels:  
            app: hello  
    - to:  
      ports:  
        - protocol: UDP  
          port: 53  
content_copy
```

1. Create an egress policy.

```
kubectl apply -f foo-allow-to-hello.yaml  
content_copy
```

2. Verify that the policy was created.

```
kubectl get networkpolicy  
content_copy
```

#### Output (Do not copy)

NAME	POD-SELECTOR	AGE
foo-allow-to-hello	app=foo	7s
hello-allow-from-foo	app=hello	5m

```
content_copy
```

## Validate the egress policy

1. Deploy a new web application called hello-web-2 and expose it internally in the cluster.

```
kubectl run hello-web-2 --labels app=hello-2 \  
  --image=gcr.io/google-samples/hello-app:1.0 --port 8080 --expose  
content_copy
```

2. Run a temporary Pod with the `app=foo` label and get a shell prompt inside the container.

```
kubectl run test-3 --labels app=foo --image=alpine --restart=Never --rm --  
  stdin --tty  
content_copy
```

3. Verify that the Pod can establish connections to hello-web:8080.

```
wget -qO- --timeout=2 http://hello-web:8080  
content_copy
```

#### Output (Do not copy)

```
If you don't see a command prompt, try pressing enter.  
/ # wget -qO- --timeout=2 http://hello-web:8080  
Hello, world!  
Version: 1.0.0  
Hostname: hello-web  
/ #  
content_copy
```

4. Verify that the Pod cannot establish connections to hello-web-2:8080.

```
 wget -qO- --timeout=2 http://hello-web-2:8080  
content_copy
```

This fails because none of the Network policies you have defined allow traffic to Pods labelled app: hello-2.

5. Verify that the Pod cannot establish connections to external websites, such as www.example.com.

```
 wget -qO- --timeout=2 http://www.example.com  
content_copy
```

This fails because the network policies do not allow external http traffic (tcp port 80).

6. Type **exit** and press **ENTER** to leave the shell.  
Click *Check my progress* to verify the objective.

Run web server applications

Check my progress

The screenshot shows the Google Cloud Platform Workloads interface. The left sidebar has categories: Clusters, Workloads (selected), Services & Ingress, and Marketplace. The main area displays a table of workloads:

Name	Status	Type	Pods	Namespace	Cluster
dns-demo-1	Running	Pod	1/1	default	standard-cluster-1
dns-demo-2	Running	Pod	1/1	default	standard-cluster-1
hello-v1	OK	Deployment	3/3	default	standard-cluster-1
hello-v2	Does not have minimum availability	Deployment	2/3	default	standard-cluster-1

At the bottom, there is a terminal window showing the command: cat hello-lb-svc.yaml

```
student_01_f77ca049039c@cloudshell:~/ak8s/GKE_Services (qwiklabs-gcp-01-ee1857955d37)$ cat hello-lb-svc.yaml
apiVersion: v1
kind: Service
metadata:
  name: hello-lb-svc
spec:
  type: LoadBalancer
  loadBalancerIP: 34.71.205.152
  selector:
    name: hello-v2
  ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
student_01_f77ca049039c@cloudshell:~/ak8s/GKE_Services (qwiklabs-gcp-01-ee1857955d37)$
```

# Creating Services and Ingress Resources

1 hourFree

Rate Lab

## Overview

In this lab, you will create two deployments of Pods and work with three different types of services, including the Ingress resource, and explore how Kubernetes services in GKE are associated with Google Cloud Network Load Balancers.

## Objectives

In this lab, you learn how to perform the following tasks:

- Observe Kubernetes DNS in action.

- Define various service types (ClusterIP, NodePort, LoadBalancer) in manifests along with label selectors to connect to existing labeled Pods and deployments, deploy those to a cluster, and test connectivity.
- Deploy an Ingress resource that connects clients to two different services based on the URL path entered.
- Verify Google Cloud network load balancer creation for type=LoadBalancer services.

## Task 0. Lab Setup

### Access Qwiklabs

For each lab, you get a new GCP project and set of resources for a fixed time at no cost.

1. Make sure you signed into Qwiklabs using an **incognito window**.

02:00:00

2. Note the lab's access time (for example, **02:00:00**) and make sure you can finish in that time block.

There is no pause feature. You can restart if needed, but you have to start at the beginning.

START LAB

3. When ready, click

4. Note your lab credentials. You will use them to sign in to Cloud Platform Console.

[Open Google Console](#)

**Caution:** When you are in the console, do not deviate from the lab instructions. Doing so may cause your account to be blocked. [Learn more.](#)

Username  
 [!\[\]\(f76933a0c8fdccd37e1cfc596589b138\_img.jpg\)](#)

Password  
 [!\[\]\(4e98a6d640bb8a5eac4cb132cc8dbc78\_img.jpg\)](#)

GCP Project ID  
 [!\[\]\(5ff33a4d86a06075466c4842a28e4f41\_img.jpg\)](#)

New to labs? [View our introductory video!](#)

5. Click **Open Google Console**.
6. Click **Use another account** and copy/paste credentials for **this** lab into the prompts.

If you use other credentials, you'll get errors or **incur charges**.

7. Accept the terms and skip the recovery resource page.

Do not click **End Lab** unless you are finished with the lab or want to restart it.

This clears your work and removes the project.

After you complete the initial sign-in steps, the project dashboard appears.

## Open Cloud Shell

You will do most of the work for this lab in [Cloud Shell](#).



1. On the Google Cloud Console title bar, click **Activate Cloud Shell** ().
2. Click **Start Cloud Shell**.

After a moment of provisioning, the Cloud Shell prompt appears:

## Task 1. Connect to the lab GKE cluster and test DNS

In this task, you connect to the lab GKE cluster and create a deployment manifest for a set of Pods within the cluster that you will then use to test DNS resolution of Pod and service names.

### Connect to the lab GKE cluster

1. In Cloud Shell, type the following command to create environment variables for the Google Cloud zone and cluster name that will be used to create the cluster for this lab.

```
export my_zone=us-central1-a  
export my_cluster=standard-cluster-1  
content_copy
```

2. Configure tab completion for the kubectl command-line tool.

```
source <(kubectl completion bash)  
content_copy
```

3. Configure access to your cluster for kubectl:

```
gcloud container clusters get-credentials $my_cluster --zone $my_zone  
content_copy
```

## Task 2. Create Pods and services to test DNS resolution

You will create a service called `dns-demo` with two sample application Pods called `dns-demo-1` and `dns-demo-2`.

### Clone the source file repository and deploy the sample application pods

The pods that you will use to test network connectivity via Kubernetes services are deployed using the `dns-demo.yaml` manifest file that has been provided for you in the source repository. You will use these pods later to test connectivity via various services from systems inside your Kubernetes cluster. You will also test connectivity from external addresses using the Cloud Shell. You can use the Cloud Shell for this because it is on a network that is completely separate to your Kubernetes cluster networks.

```
apiVersion: v1
kind: Service
metadata:
  name: dns-demo
spec:
  selector:
    name: dns-demo
  clusterIP: None
  ports:
    - name: dns-demo
      port: 1234
      targetPort: 1234
---
apiVersion: v1
kind: Pod
metadata:
  name: dns-demo-1
  labels:
    name: dns-demo
spec:
  hostname: dns-demo-1
  subdomain: dns-demo
  containers:
    - name: nginx
      image: nginx
---
apiVersion: v1
```

```
kind: Pod
metadata:
  name: dns-demo-2
  labels:
    name: dns-demo
spec:
  hostname: dns-demo-2
  subdomain: dns-demo
  containers:
  - name: nginx
    image: nginx
content_copy
```

1. In Cloud Shell enter the following command to clone the repository to the lab Cloud Shell.

```
git clone https://github.com/GoogleCloudPlatform/training-data-analyst
content_copy
```

2. Create a soft link as a shortcut to the working directory.

```
ln -s ~/training-data-analyst/courses/ak8s/v1.1 ~/ak8s
content_copy
```

3. Change to the directory that contains the sample files for this lab.

```
cd ~/ak8s/GKE_Services/
content_copy
```

4. Create the service and Pods.

```
kubectl apply -f dns-demo.yaml
content_copy
```

5. Verify that your Pods are running.

```
kubectl get pods
content_copy
```

The output should look like this example.

## Output (do not copy)

NAME	READY	STATUS	RESTARTS	AGE
dns-demo-1	1/1	Running	0	8s
dns-demo-2	1/1	Running	0	8s

```
content_copy
```

Click *Check my progress* to verify the objective.

Create Pods and services to test DNS resolution

Check my progress

6. To get inside the cluster, open an interactive session to bash running from dns-demo-1.

```
kubectl exec -it dns-demo-1 -- /bin/bash  
content_copy
```

Now that you are inside a container in the cluster, subsequent commands will run from that context and you will use this to test network connectivity by pinging various targets in later tasks. However, you don't have a tool to ping in this container, so you need to install the `ping` command first.

7. Update `apt-get` and install a `ping` tool.

```
apt-get update  
apt-get install -y iputils-ping  
content_copy
```

8. Ping `dns-demo-2`:

```
ping dns-demo-2.dns-demo.default.svc.cluster.local  
content_copy
```

This ping should succeed and report that the target has the ip-address you found earlier for the `dns-demo-2` Pod.

9. Press `Ctrl+C` to abort the ping command.

10. Ping the `dns-demo` service's FQDN, instead of a specific Pod inside the service:

```
ping dns-demo.default.svc.cluster.local  
content_copy
```

This ping should also succeed but it will return a response from the FQDN of one of the two demo-dns Pods. This Pod might be either `demo-dns-1` or `demo-dns-2`.

11. Press `Ctrl+C` to abort the ping command.

When you deploy applications, your application code runs inside a container in the cluster, and thus your code can access other services by using the FQDNs of those services from inside that cluster. This approach is simpler than using IP addresses or even Pod names, because those are more likely to change.

12. Leave the interactive shell on `dns-demo-1` running.

# Task 3. Deploy a sample workload and a ClusterIP service

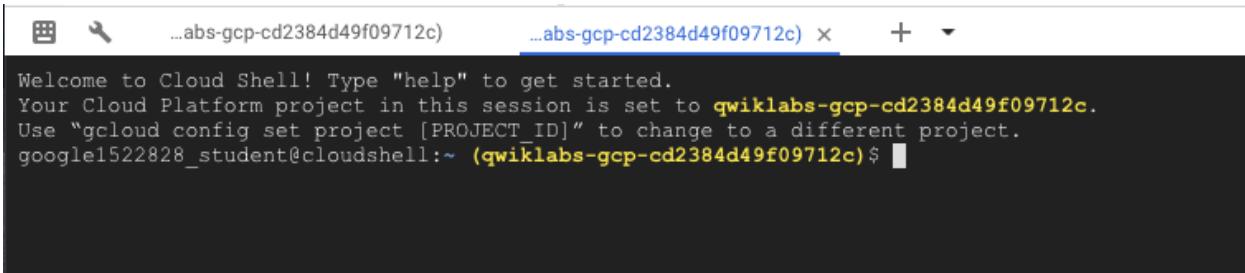
In this task, you create a deployment manifest for a set of Pods within the cluster and then expose them using a ClusterIP service.

## Deploy a sample web application to your Kubernetes Engine cluster

In this task you will deploy a sample web application container image that listens on an HTTP server on port 8080 using the following manifest that has already been created for you in the `hello-v1.yaml` manifest file.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-v1
spec:
  replicas: 3
  selector:
    matchLabels:
      run: hello-v1
  template:
    metadata:
      labels:
        run: hello-v1
        name: hello-v1
    spec:
      containers:
        - image: gcr.io/google-samples/hello-app:1.0
          name: hello-v1
          ports:
            - containerPort: 8080
              protocol: TCP
content_copy
```

1. In the Cloud Shell menu bar, click the plus sign (+) icon to start a new Cloud Shell session.



Welcome to Cloud Shell! Type "help" to get started.  
Your Cloud Platform project in this session is set to **qwiklabs-gcp-cd2384d49f09712c**.  
Use "gcloud config set project [PROJECT ID]" to change to a different project.  
google1522828\_student@cloudshell:~ (qwiklabs-gcp-cd2384d49f09712c) \$

A second Cloud Shell session appears in your Cloud Shell window. You can switch sessions by clicking them in the menu bar.

2. In the second Cloud Shell tab, change to the directory that contains the sample files for this lab.

```
cd ~/ak8s/GKE_Services/  
content_copy
```

3. To create a deployment from the `hello-v1.yaml` file, execute the following command:

```
kubectl create -f hello-v1.yaml  
content_copy
```

4. To view a list of deployments, execute the following command:

```
kubectl get deployments  
content_copy
```

The output should look like this example.

### Output (do not copy)

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-v1	3/3	3	3	14s

```
content_copy
```

## Define service types in the manifest

In this task you will deploy a Service using a ClusterIP using the `hello-svc.yaml` sample manifest that has already been created for you.

```
apiVersion: v1  
kind: Service  
metadata:
```

```
name: hello-svc
spec:
  type: ClusterIP
  selector:
    name: hello-v1
  ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
content_copy
```

1. To deploy the ClusterIP service, execute the following command:

```
kubectl apply -f ./hello-svc.yaml
content_copy
```

This manifest defines a ClusterIP service and applies it to Pods that correspond to the selector. In this case, the manifest is applied to the hello-v1 Pods that you deployed. This service will automatically be applied to any other deployments with the `name: hello-v1` label.

2. Verify that the service was created and that a Cluster-IP was allocated:

```
kubectl get service hello-svc
content_copy
```

Your IP address might be different from the example output.

## Output (do not copy)

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-svc	ClusterIP	10.11.253.203	<none>	80/TCP	20s

```
content_copy
```

No external IP is allocated for this service. Because the Kubernetes Cluster IP addresses are not externally accessible by default, creating this service does not make your application accessible outside of the cluster.

Click *Check my progress* to verify the objective.

Deploy a sample workload and a ClusterIP service

[Check my progress](#)

## Test your application

1. In the Cloud Shell, attempt to open an HTTP session to the new service using the following command:

```
curl hello-svc.default.svc.cluster.local  
content_copy
```

The connection should fail because that service is not exposed outside of the cluster.

### Output (do not copy)

```
curl: (6) Could not resolve host: hello-svc.default.svc.cluster.local  
content_copy
```

Now you will test the service from inside the cluster using the interactive shell you have running on the `dns-demo-1` Pod.

2. Return to your first Cloud Shell window, which is currently redirecting the `stdin` and `stdout` of the `dns-demo-1` Pod.
3. Install `curl` so you can make calls to web services from the command line.

```
apt-get install -y curl  
content_copy
```

4. Use the following command to test the HTTP connection between the Pods.

```
curl hello-svc.default.svc.cluster.local  
content_copy
```

This connection should succeed and provide a response similar to the output below. Your hostname might be different from the example output.

### Output (do not copy)

```
root@dns-demo-1:/# curl hello-svc.default.svc.cluster.local  
Hello, world!  
Version: 1.0.0  
Hostname: hello-v1-85b99869f-mp4vd  
root@dns-demo-1:/#  
content_copy
```

This connection works because the clusterIP can be resolved using the internal DNS within the Kubernetes Engine cluster.

## Task 4. Convert the service to use NodePort

In this task, you will convert your existing ClusterIP service to a NodePort service and then retest access to the service from inside and outside the cluster.

A modified version of the `hello-svc.yaml` file called `hello-nodeport-svc.yaml` that changes the service type to NodePort has already been created for you.

```
apiVersion: v1
kind: Service
metadata:
  name: hello-svc
spec:
  type: NodePort
  selector:
    name: hello-v1
  ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
    nodePort: 30100
content_copy
```

1. Return to your second Cloud Shell window. This is the window NOT connected to the stdin and stdout of the `dns-test` Pod.
2. To deploy the manifest that changes the service type for the `hello-svc` to NodePort , execute the following command:

```
kubectl apply -f ./hello-nodeport-svc.yaml
content_copy
```

This manifest redefines `hello-svc` as a NodePort service and assigns the service port 30100 on each node of the cluster for that service.

3. Enter the following command to verify that the service type has changed to NodePort:

```
kubectl get service hello-svc
content_copy
```

Your IP address might be different from the example output.

**Output (do not copy)**

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
hello-svc	NodePort	10.11.253.203	<none>	80:30100/TCP	59m
<u>content_copy</u>					

Note that there is still no external IP allocated for this service.

Click *Check my progress* to verify the objective.

Convert the service to use NodePort

[Check my progress](#)

## Test your application

1. In the Cloud Shell, attempt to open an HTTP session to the new service using the following command:

```
curl hello-svc.default.svc.cluster.local  
content_copy
```

The connection should fail because that service is not exposed outside of the cluster.

### Output (do not copy)

```
curl: (6) Could not resolve host: hello-svc.default.svc.cluster.local  
content_copy
```

Now you will test the service from another Pod.

2. Return to your first Cloud Shell window, which is currently redirecting the stdin and stdout of the dns-test Pod.
3. Use the following command to test the HTTP connection between the Pods.

```
curl hello-svc.default.svc.cluster.local  
content_copy
```

This connection should succeed and provide a response similar to the output below. Your hostname might be different from the example output.

### Output (do not copy)

```
root@dns-demo-1:/# curl hello-svc.default.svc.cluster.local
```

```
Hello, world!
Version: 1.0.0
Hostname: hello-v1-85b99869f-mp4vd
root@dns-demo-1:/#
content_copy
```

This connection works because the clusterIP can be resolved using the internal DNS within the GKE cluster.

## Task 5. Create static public IP addresses using Google Cloud Networking

### Reserve Google Cloud Networking Static IP Addresses

1. In the Google Cloud Console navigation menu, navigate to **Networking > VPC Network > External IP Addresses**.
2. Click **+ Reserve Static Address**.
3. Enter `regional-loadbalancer` for the **Name**.
4. Explore the options but leave the remaining settings at their defaults. Note that the default type is **Regional**.

Note that the default region must match the region where you have deployed your GKE cluster. If this is not the case then change the region here to match the region you used to create your cluster. If you accepted the lab defaults then that should be the us-central1 (Iowa) region.

5. Click **Reserve**.

6. Make a note of the external ip-address called `regional-loadbalancer`. You will use this in a later task.
7. Click **+ Reserve Static Address**.
8. Enter `global-ingress` for the **Name**.
9. Change the Type to **Global**.
10. Leave the remaining settings at their defaults.
11. Click **Reserve**.
12. Make a note of the external ip-address called `global-ingress`. You will use this in a later task.

Click *Check my progress* to verify the objective.

Create static public IP addresses using Google Cloud Networking

[Check my progress](#)

## Task 6. Deploy a new set of Pods and a LoadBalancer service

You will now deploy a new set of Pods running a different version of the application so that you can easily differentiate the two services. You will then expose the new Pods as a LoadBalancer service and access the service from outside the cluster.

A second Deployment manifest called `hello-v2.yaml` has been created for you that creates a new deployment that runs version 2 of the sample hello application on port 8080.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: hello-v2
spec:
  replicas: 3
  selector:
    matchLabels:
      run: hello-v2
  template:
    metadata:
      labels:
        run: hello-v2
        name: hello-v2
    spec:
      containers:
        - image: gcr.io/google-samples/hello-app:2.0
          name: hello-v2
          ports:
            - containerPort: 8080
              protocol: TCP
content_copy
```

1. Return to your second Cloud Shell window. This is the window NOT connected to the stdin and stdout of the dns-test Pod.
2. To deploy the manifest that creates the `hello-v2` deployment execute the following command:

```
kubectl create -f hello-v2.yaml
content_copy
```

3. To view a list of deployments, execute the following command:

```
kubectl get deployments
content_copy
```

The output should look like this example.

## Output (do not copy)

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
hello-v1	3/3	3	3	18m
hello-v2	3/3	3	3	7s

```
content_copy
```

## Define service types in the manifest

In this task you will deploy a LoadBalancer Service using the `hello-lb-svc.yaml` sample manifest that has already been created for you.

You will use the `sed` command to replace the `10.10.10.10` placeholder address in the load balancer yaml file with the static address you reserved earlier for the load balancer.

```
apiVersion: v1
kind: Service
metadata:
  name: hello-lb-svc
spec:
  type: LoadBalancer
  loadBalancerIP: 10.10.10.10
  selector:
    name: hello-v2
  ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
content_copy
```

1. Make sure you are still in the second Cloud Shell window. This is the window NOT connected to the `stdin` and `stdout` of the `dns-test` Pod.
2. Enter the following command in the Cloud Shell to save the regional static IP-address you created earlier into an environment variable.

```
export STATIC_LB=$(gcloud compute addresses describe regional-loadbalancer --region us-central1 --format json | jq -r '.address')content_copy
```

3. Enter the following command in the Cloud Shell to replace the placeholder address with the regional static IP-address.

```
sed -i "s/10\.10\.10\.10/$STATIC_LB/g" hello-lb-svc.yaml
content_copy
```

4. To check that the external ip-address has been replaced correctly enter the following command in the Cloud shell.

```
cat hello-lb-svc.yaml
content_copy
```

The `loadBalancerIP` should match the address you recorded earlier for the `regional-loadbalancer` reserved static ip-address.

- To deploy your LoadBalancer service manifest, execute the following command:

```
kubectl apply -f ./hello-lb-svc.yaml  
content_copy
```

This manifest defines a LoadBalancer service, which deploys a Google Cloud Network Load Balancer to provide external access to the service. This service is only applied to the Pods with the `name: hello-v2` selector.

- Verify that the service was created and that a node port was allocated:

```
kubectl get services  
content_copy
```

Your IP address might be different from the example output.

### Output (do not copy)

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
dns-demo	ClusterIP	None	<none>	1234/TCP	1h
hello-lb-svc	LoadBalancer	10.11.253.48	35.184.45.240	80:30103/TCP	2m
hello-svc	NodePort	10.11.253.203	<none>	80:30100/TCP	1h
kubernetes	ClusterIP	10.11.240.1	<none>	443/TCP	2h

Notice that the new LoadBalancer service lists the `regional-loadbalancer` reserved address as its external IP address. In GKE environments the external load balancer for the LoadBalancer service type is implemented using a Google Cloud load balancer and will take a few minutes to create. This external IP address makes the service accessible from outside the cluster.

## Confirm that a Regional Google Cloud Load Balancer has been created

- In the Google Cloud Console Navigate to **Networking > Network Services > Load Balancing**
- You should see a TCP Load Balancer listed with 1 target pool backend and 2 instances. Click the name to open the details page.

3. The details should show that it is a regional TCP load balancer using the static IP address you called `regional-loadbalancer` when you reserved it earlier.

Click *Check my progress* to verify the objective.

Deploy a new set of Pods and a LoadBalancer service

Check my progress

## Test your application

1. In the Cloud Shell, attempt to open an HTTP session to the new service using the following command:

```
curl hello-lb-svc.default.svc.cluster.local  
content_copy
```

The connection should fail because that service name is not exposed outside of the cluster.

### Output (do not copy)

```
curl: (6) Could not resolve host: hello-lb-svc.default.svc.cluster.local  
content_copy
```

This occurs because the external IP address is not registered with this hostname..

2. Try the connection again using the External IP address associated with the REgional Load Balancer service. Type the following command and substitute `[external_IP]` with your service's external IP address.

```
curl [external_IP]  
content_copy
```

### Example (do not copy)

```
curl 35.184.45.240  
content_copy
```

This time the connection does not fail because the LoadBalancer's external ip-address can be reached from outside Google Cloud.

**Note:** This may take up to 5 minutes to become available.

## Output (do not copy)

```
Hello, world!
Version: 2.0.0
Hostname: hello-v2-59c99d8b65-jrx2d
content_copy
```

3. Return to your first Cloud Shell window, which is currently redirecting the stdin and stdout of the `dns-demo-1` Pod.
4. Use the following command to test the HTTP connection between the Pods.

```
curl hello-lb-svc.default.svc.cluster.local
content_copy
```

This connection should succeed and provide a response similar to the output below. Your hostname will be different from the example output.

## Output (do not copy)

```
root@dns-demo-1:/# curl hello-lb-svc.default.svc.cluster.local
Hello, world!
Version: 2.0.0
Hostname: hello-v2-59c99d8b65-78ggz
root@dns-demo-1:/#
content_copy
```

The internal DNS name works within the Pod, and you can see that you are accessing the same v2 version of the application as you were from outside of the cluster using the external ip address.

5. Try the connection again within the Pod using the External IP address associated with the service. Type the following command and substitute `[external_IP]` with your service's external IP address.

```
curl [external_IP]
content_copy
```

## Example (do not copy)

```
curl 35.184.45.240
content_copy
```

## Output (do not copy)

```
root@dns-demo-1:/# curl 35.184.45.240
Hello, world!
Version: 2.0.0
Hostname: hello-v2-59c99d8b65-jrx2d
root@dns-demo-1:/#
```

```
content_copy
```

The external IP also works from inside Pods running in the cluster, and returns a result from the same v2 version of the applications.

4. Exit the console redirection session to the Pod by typing exit.

```
exit
```

```
content_copy
```

This will return you to the Cloud Shell command prompt.

## Task 7. Deploy an Ingress resource

You have two services in your cluster for the hello application. One service is hosting version 1.0 via a NodePort service, while the other service is hosting version 2.0 via a LoadBalancer service. You will now deploy an Ingress resource that will direct traffic to both services based on the URL entered by the user.

### Create an ingress resource

Ingress is a Kubernetes resource that encapsulates a collection of rules and configuration for routing external HTTP(S) traffic to internal services.

On GKE, Ingress is implemented using Cloud Load Balancing. When you create an ingress resource in your cluster, GKE creates an HTTP(S) load balancer and configures it to route traffic to your application.

A sample manifest called `hello-ingress.yaml` that will configure an ingress resource has been created for you.

```
apiVersion: app/v1
kind: Ingress
metadata:
```

```
name: hello-ingress
annotations:
  nginx.ingress.kubernetes.io/rewrite-target: /
  kubernetes.io/ingress.global-static-ip-name: "global-ingress"

spec:
  rules:
  - http:
    Paths:
    - path: /v1
      backend:
        serviceName: hello-svc
        servicePort: 80
    - path: /v2
      backend:
        serviceName: hello-lb-svc
        servicePort: 80
```

content\_copy

This configuration file defines an ingress resource that will attach to the `global-ingress` static ip-address you created in a previous step. This ingress resource will create a global HTTP(S) load balancer that directs traffic to your web services based on the path entered.

The `kubernetes.io/ingress.global-static-ip-name` annotation allows you to specify a named reserved ip-address that the Ingress resource will use when creating the Google Cloud Global HTTP(S) load balancer for the Ingress resource.

1. To deploy this ingress resource, execute the following command:

```
kubectl apply -f hello-ingress.yaml
```

content\_copy

When you deploy this manifest, Kubernetes creates an ingress resource on your cluster. The ingress controller running in your cluster is responsible for creating an HTTP(S) load balancer to route all external HTTP traffic (on port 80) to the web NodePort service and the LoadBalancer service that you exposed.

## Confirm that a Global HTTP(S) Load Balancer has been created

1. In the Google Cloud Console Navigate to **Networking > Network Services > Load Balancing**

2. You should now also see a HTTP(S) Load Balancer listed. Click the name to open the details page.
3. The details should show that it is a global HTTP(S) load balancer using the static IP address you called `global-ingress` when you reserved it earlier.

Click *Check my progress* to verify the objective.

Deploy an Ingress resource

Check my progress

## Test your application

1. To get the external IP address of the load balancer serving your application, execute the following command:

```
kubectl describe ingress hello-ingress  
content_copy
```

**Partial output (do not copy)**

```
Name:          hello-ingress
Namespace:    default
Address:      35.244.173.44
Default backend: default-http-backend:80 (10.8.2.5:8080)
Rules:
  Host    Path  Backends
  ----  ----  -----
  *
    /v1    hello-svc:80 (<none>)
    /v2    hello-lb-svc:80 (<none>)
[...]
  ingress.kubernetes.io/backends: {"k8s-[*]": "HEALTHY", [*]}
[...]

Events:
  Type  Reason  Age From                Message
  ----  -----  --  ----                -----
  Normal ADD     5m  loadbalancer-controller  default/hello-ingress
  Normal CREATE   5m  loadbalancer-controller ip: 35.244.173.44
content_copy
```

**Note:** You need to wait for a few minutes for the load balancer to become active, and for the health checks to succeed, before the external address will be displayed.

Repeat the command every few minutes to check if the Ingress resource has finished initializing. When the Google Cloud global HTTP(S) load balancer has

been created and initialized the command will report the external ip-address which will match the global static ip-address you reserved earlier that you called `global-ingress`.

2. Use the External IP address associated with the Ingress resource, and type the following command, substituting `[external_IP]` with the Ingress resource's external IP address. Be sure to include the `/v1` in the URL path.

```
curl http://[external_IP]/v1  
content_copy
```

#### Example (do not copy)

```
curl http://35.201.92.69/v1  
content_copy
```

#### Output (do not copy)

```
Hello, world!  
Version: 1.0.0  
Hostname: hello-v1-85b99869f-4rmqw  
content_copy
```

The `v1` URL is configured in `hello-ingress.yaml` to point to the `hello-svc` NodePort service that directs traffic to the `v1` application Pods..

**Note:** GKE might take a few minutes to set up forwarding rules until the Global load balancer used for the Ingress resource is ready to serve your application. In the meantime, you might get errors such as HTTP 404 or HTTP 500 until the load balancer configuration is propagated across the globe.

3. Now test the `v2` URL path from Cloud Shell. Use the External IP address associated with the Ingress resource, and type the following command, substituting `[external_IP]` with the Ingress resource's external IP address. Be sure to include the `/v2` in the URL path.

```
curl http://[external_IP]/v2  
content_copy
```

The `v2` URL is configured in `hello-ingress.yaml` to point to the `hello-lb-svc` LoadBalancer service that directs traffic to the `v2` application Pods.

#### Example (do not copy)

```
curl http://35.201.92.69/v2  
content_copy
```

#### Output (do not copy)

```
Hello, world!
```

```
Version: 2.0.0
Hostname: hello-v2-59c99d8b65-jrx2d
content_copy
```

## Inspect the changes to your networking resources in the Google Cloud Console

1. In Google Cloud Console, on the **Navigation menu** , click **Networking > Network services> Load balancing**.

There are now two load balancers listed:

There is the initial regional load balancer for the `hello-lb-svc` service. This has a UID style name and is configured to load balance TCP port 80 traffic to the cluster nodes.

The second was created for the Ingress object and is a full HTTP(S) load balancer that includes host and path rules that match the Ingress configuration. This will have `hello-ingress` in its name.

2. Click the load balancer with `hello-ingress` in the name.

This will display the summary information about the protocols, ports, paths and backend services of the Google Cloud Global HTTP(S) load balancer created for your Ingress resource.

Quiz: Google Kubernetes Engine Networking

Your score: 83% Passing score: 83%

Congratulations! You passed this assessment.

Retake

1. You have updated your application and deployed a new Pod. What step can you take to ensure that you can access the Pod and the Pod's application throughout the lifecycle of the Pod using a durable IP address?

- Add the fully qualified domain name of the application's Pod to your local hostfile.
- Deploy a Kubernetes Service with a selector that locates the application's Pods.
- Register the fully qualified domain name of the application's Pod in DNS.
- Add metadata annotations to the Pod manifest that define a persistent DNS name.

That is correct.

Chat

Quiz: Google Kubernetes Engine Networking

3. Your Pod has been rescheduled and the IP address that was assigned to the Pod when it was originally scheduled is no longer accessible. What is the reason for this?

- The old Pod IP address is blocked by a firewall.
- The new Pod has received a different IP address.
- The new Pod IP address is blocked by a firewall.
- The Pod IP range for the cluster is exhausted.

That is correct.

4. You need to apply a network policy to five Pods to block ingress traffic from other Pods. Each Pod has a label of app:demo-app. In your network policy manifest, you have specified the label app:demo-app in spec.selector. The policy is configured and when you list the Network

Chat

Quiz: Google Kubernetes Engine Networking

4. You need to apply a network policy to five Pods to block ingress traffic from other Pods. Each Pod has a label of app:demo-app. In your network policy manifest, you have specified the label app:demo-app in spec.selector. The policy is configured and when you list the Network Policies on your cluster you can see the policy listed but is not having any effect as you can still ping the Pods from other Pods in the cluster. What is the problem and what action can you take to correct this?

- Network policies are only applied when a Pod is started. Stop all of the Pods in your application and restart them in order to activate the network policy.
- You have not enabled network policies on the cluster. Enable network policies on the cluster, and reboot all of the nodes.
- A network policy must be applied to all Pods in the cluster in order to block ingress traffic. In the network policy manifest, do not define any value for spec.selector.
- You need to create a matching Google Cloud firewall rule for the network policy. In the Cloud Console or Cloud Shell create a firewall rule that matches the Network Policy.

Chat

Quiz: Google Kubernetes Engine Networking

5. What change can an administrator make to achieve the lowest possible latency when using a Google Cloud load-balancer service to balance network traffic, minimizing double-hop traffic between nodes? Assume that container-native load balancing is not in use.

- Set the externalTrafficPolicy field to roundRobin in the YAML manifest.
- Set the externalTrafficPolicy field to local in the YAML manifest.
- Set the spec.type field to LoadBalancer in the YAML manifest.
- Set the spec.type field to Local in the YAML manifest.

That is correct.

6. During testing you cannot find the Google Cloud Load Balancer that should have been configured for your application. You check the manifest for the application and notice that the application's front-end Service type is ClusterIP. How can you correct this?

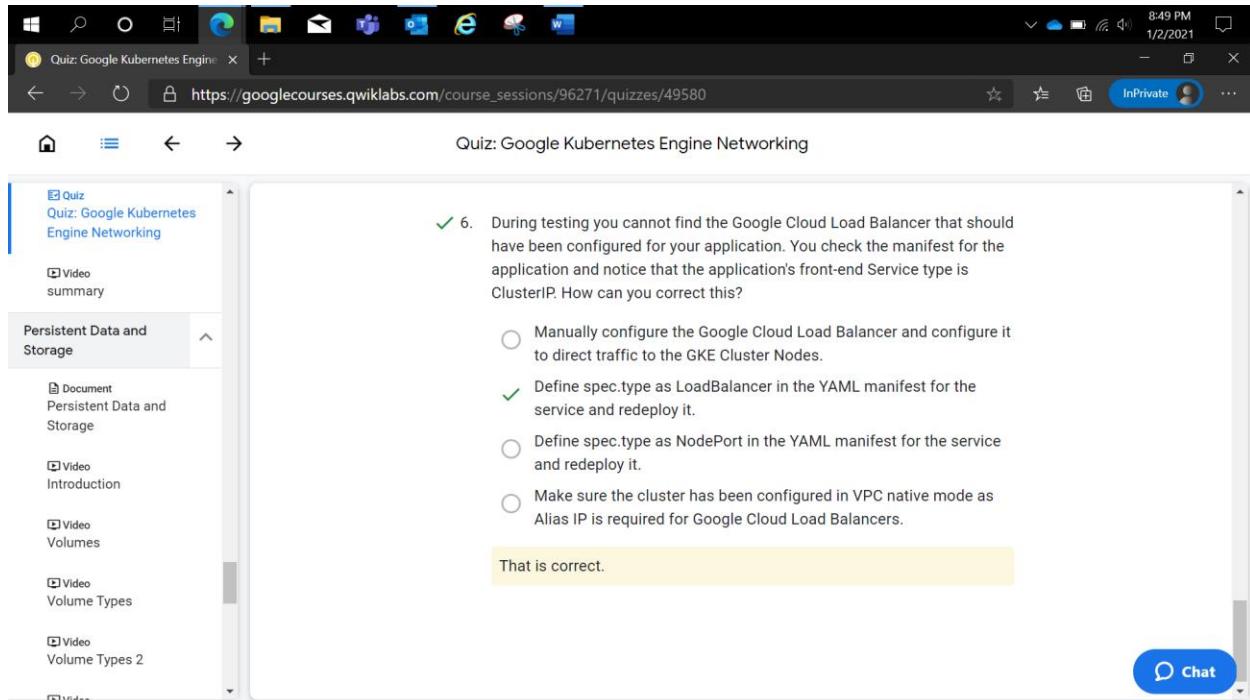
Chat

Quiz: Google Kubernetes Engine Networking

6. During testing you cannot find the Google Cloud Load Balancer that should have been configured for your application. You check the manifest for the application and notice that the application's front-end Service type is ClusterIP. How can you correct this?

- Manually configure the Google Cloud Load Balancer and configure it to direct traffic to the GKE Cluster Nodes.
- Define spec.type as LoadBalancer in the YAML manifest for the service and redeploy it.
- Define spec.type as NodePort in the YAML manifest for the service and redeploy it.
- Make sure the cluster has been configured in VPC native mode as Alias IP is required for Google Cloud Load Balancers.

That is correct.



Quiz: Google Kubernetes Engine Networking

6. You are designing a GKE solution. One of your requirements is that network traffic load balancing should be directed to Pods directly, instead of balanced across nodes. What change can you make to your environment?

- Configure external access using the Nodeport Service type working with an external network load balancer.
- Configure all external access for your application using Ingress resources rather than Services.
- Configure or migrate your cluster to VPC-Native Mode and deploy a Container-native load balancer.
- Set the externalTrafficPolicy field to local in the YAML manifest for your external services.
- Configure affinity and antiaffinity rules that ensure your application's Pods are distributed across Nodes.

That is correct.

